

Day 1 — RAG Foundations (From First Principles)

Goal

Understand what Retrieval-Augmented Generation (RAG) is and why naive prompting fails.

Day 1 establishes the mental model: answers must come from evidence, not the model's memory.

Core Idea

RAG separates *knowledge retrieval* from *language generation*.

The LLM is treated as a reasoning engine, not a database.

Flow

User Query → Document Corpus → Retrieval → Context → Answer

At this stage, there is no safety or validation—only plumbing.

Key Takeaway

Without structure and constraints, LLMs will hallucinate confidently.

RAG is necessary but not sufficient.